

582631 Introduction to Machine Learning

Period II, Autumn 2015

Exercise #1 (Python for Machine Learning)

Tuesday, 27 October, 12:15-16 in B221

(No preparation necessary, no handing-in of solutions, no points for exercises.)

Exercise 1. Generate 100 samples from a normal distribution with zero mean and unit variance.

- (a) Visualize the data by drawing a histogram.
- (b) Sort the data points into increasing order.
- (c) Draw a line plot of points (x_i, y_i) for $i = 1 \dots 100$ where x_i is the i :th data point (in the sorted order) and y_i is equal to $i/100$.
- (d) Can you tell what you have drawn?

Numpy: random, sort, arange;

matplotlib.pyplot: plot;

Exercise 2. Generate a random matrix (dimensions for example 10×5) and store it into a file. Write a function that

- Takes a filename as an argument.
- Reads a data matrix X from the given file.
- Outputs row and column sums of X as bar plots. Try to put both plots into the same window.
- Returns two values: the dimensions of the matrix as a vector and the sum of all elements.

Test your function by calling the function with the name of the file.

Numpy: save, load, sum;

matplotlib.pyplot: subplots, bar;

Exercise 3. The law of large numbers states that the average of a series of i.i.d. samples approaches the expectation of the generating distribution as the number of samples goes to infinity. More formally, if all x_n are i.i.d. sampled from some random variable X we have

$$\lim_{N \rightarrow \infty} \frac{1}{N} \sum_{n=1}^N x_n = E(X)$$

with probability 1. Let's verify this theorem empirically.

- (a) Generate samples from a normal distribution with mean 0 and variance 1.
- (b) Plot N vs. the empirical mean $\frac{1}{N} \sum_{n=1}^N x_n$ in logarithmic scale.
- (c) How fast does the empirical mean approach the expectation?

Numpy: random, arange, sum, log;

matplotlib.pyplot: plot;

Exercise 4. The expectation of a random variable is a linear operator:

$$\begin{aligned}E(\alpha X) &= \alpha E(X) \\E(X + Y) &= E(X) + E(Y)\end{aligned}$$

Let F^n be a random variable describing the number of fixed points (elements that don't change place) in a random permutation of n elements.

- (a) Write a function that takes n as an argument and returns a sample of length k from the distribution of F^n .
- (b) Generate samples from F^n with a few different values of n .
- (c) Draw histograms of the different samples.
- (d) Can you guess what the expected value $E(F^n)$ of the distribution might be?
- (e) Use the linearity of the expectation operator to calculate $E(F^n)$. (Hint: Take another random variable F_i^n , such that $F_i^n = 1$ when the i :th element stays fixed, otherwise $F_i^n = 0$. Now we can write $F^n = \sum_{i=1}^n F_i^n$.)

Numpy: permutation;

matplotlib.pyplot: hist;

Exercise 5.

- (a) Write a function that
 - takes 6 arguments: a sample size n , a probability π ($0 \leq \pi \leq 1$), and the expectatations and variances of two univariate normal distributions.
 - returns two n -element vectors \mathbf{x} and \mathbf{y} , such that for each $i \in \{1, \dots, n\}$ holds:
With probability π , $\mathbf{y}[i] = 0$ and $\mathbf{x}[i]$ is a point from the first normal distribution
otherwise $\mathbf{y}[i] = 1$, and $\mathbf{x}[i]$ is a point from the second normal distribution.
(The distribution of \mathbf{x} is called a “Gaussian mixture”.)
- (b) Visualize the data generated by the function with different values of the arguments.
- (c) Can you come up with any example that follows such a distribution?

Numpy: random;

matplotlib.pyplot: scatter/hist;